

Sample Paper for the Erti Harness

Abstract. Attention-only models replaced recurrence in machine translation because they train in parallel and reach better quality at lower cost. This short document exists so that tests can open, read and mark up a real PDF without shipping anyone else's work.

1 Introduction

Self-attention relates every position of a sequence to every other position in a single step. Multi-head attention runs several of these in parallel, each free to attend to a different kind of relation.

2 Method

We employ eight parallel attention layers, or heads. Each head uses a smaller dimension, so the total cost is close to one full-size head.

3 Results

Pre-trained bidirectional encoders extended the approach to a wide range of tasks, reporting strong results on eleven benchmarks.

4 Limitations

Attention is quadratic in sequence length, which limits long inputs.